

Automating the Management of Software Systems

DRAFT [in progress]

Background	3
Design Semantics and Behavioral Correctness	3
Semantic Invariants	4
Environmental Conditions and Learned Invariants	4
Design Driven Automation	5
Design Instrumentation and Semantic Transformation	5
Lossless Behavioral Representation and Domain-Specific Compression	6
The Design Feedback Loop	7
The Design Learning Platform	7

Background

Modern software systems are vast, often spanning thousands of interconnected services and components distributed across data centers and geographic regions. These systems process enormous volumes of requests and evolve continuously through frequent deployments. In such environments, failures or disruptions are costly and degrade the user experience, necessitating the development of intelligent tools that can automatically manage software systems to ensure their reliability, performance and resilience.

Traditionally, to understand the execution of a software system, the engineer relied on log files to provide clues about what the system was doing during execution. Typically, this was in the form of unstructured logs and it could include a stack trace provided by an exception. These clues would then be used to understand the behavior of the design that led to the unwanted state. In some cases, the bug is the result of an incorrect implementation of the design and in other cases, the design learns about constraints imposed by the environment that it was unaware of. While this is already a tedious process, the distributed nature of modern software systems adds to the challenge. There have been many interesting approaches to expedite the debugging and recovery process by understanding the behavior of the design from the logs but all these solutions inherently suffer from the same limitation, they are working with incomplete information and as a result, their conclusions are inherently probabilistic instead of verifiable facts.

To address these limitations, this document proposes an approach that grounds the understanding of execution in a faithful, semantic representation of the system's design. This representation enables execution to be interpreted deterministically, providing a foundation for automated analysis and management without reliance on probabilistic inference.

Design Semantics and Behavioral Correctness

Before exploring how these limitations can be addressed, it is useful to clarify how the execution of a software system is understood through the behavior of its design. In this context, behavior is defined as the actions taken and choices made by the design in response to a given input. The design unambiguously establishes the intentions of a world governed by semantic contracts and relationships, expressed as semantic invariants that define the conditions under which the

world is valid. When the design's behavior preserves these invariants, the world functions as intended; when an invariant is violated, the world enters a semantically invalid state.

Semantic Invariants

These semantic invariants constrain the design's behavior to ensure that its intentions are realized. Because invariants define the conditions that must always hold true for the design to function correctly, they make failure modes predictable. When a potential failure is identified, it is resolved by refining the design's behavior to enforce the violated invariant.

For example, consider a library manager whose design accepts books and places them on shelves using the first letter of each book's name as a sorting key. Within this design, accepting a book without a name violates a semantic invariant and predicts a failure when the behavior attempts to read the first letter. The resolution is a behavioral refinement of the design: rejecting books without a name at the point of acceptance. By enforcing this invariant, the design prevents the system from entering a semantically invalid state.

Environmental Conditions and Learned Invariants

However, the assumptions made by a design can conflict with the reality of what is actually possible. A design encodes expectations about its operating environment, and those expectations may be refuted by observed behavior during execution. When this occurs, the environment imposes restrictions on how the design can realize its intentions, forcing the design to learn new behavior and introduce new constraints.

For example, in the library manager, the design assumes that it can place all the books in a basket onto a shelf in a single operation. In practice, if the basket contains more than five books, it cannot be carried to the bookshelf because it becomes too heavy. The environment therefore introduces a previously unknown constraint. To continue realizing its intention of placing books on the shelf, the design must adapt its behavior by only carrying up to five books at a time.

Through this process, the design acquires new, evidence-backed constraints derived from execution. These constraints both guide future behavior and serve as provable conditions under which the design's intentions can be realized, transforming environmental restrictions into explicit knowledge embedded within the design itself.

Design Driven Automation

In software systems, a design is realized through an implementation in a programming language, which uses its abstractions to enact the design's intentions while enforcing the semantic invariants that define correctness. When the implementation is executed and the system enters a semantically invalid state, debugging can be understood as interpreting the execution through the behavior of the design in order to restore alignment with its intentions. Because semantic invariants are unambiguous, the root cause of an invalid state can be identified by determining how the implementation violated one or more invariants. If no known invariant violation can be identified, the invalid state instead reveals a limitation imposed by the environment, enabling the design to learn and incorporate new invariants or behavioral refinements. This process eliminates probabilistic diagnosis, every semantically invalid state is either explained by an invariant violation or results in new semantic knowledge being incorporated into the design.

To address the limitations of traditional approaches to understanding software systems, this document presents an approach that automatically interprets execution through a faithful representation of the design's behavior. Because execution is understood in terms of semantic invariants, this approach has several direct implications.

- First, it automates debugging by deterministically identifying which semantic invariants were violated when the system enters a semantically invalid state.
- Second, it automates testing, since an implementation that preserves all semantic invariants is, by definition, capable of realizing the design's intentions.
- Third, it automates systems management by continuously ensuring that execution remains aligned with the design's intentions as the system operates within its environment.

Design Instrumentation and Semantic Transformation

To make this possible, the design is unambiguously instrumented into the implementation using a Design Abstraction Language (DAL), such that the abstractions of the design are mapped directly onto the implementation. As a result, the execution of the software system can be transformed into the behavior of the design, allowing execution to be automatically understood

through the design's semantics. This transformation is referred to as a Semantic Transform (ST).

For designs that involve concurrency, the Semantic Transform maps execution distributed across multiple threads, processes, or nodes into the behavior of a single, coherent design. This is achieved by propagating unique identifiers across threads, processes, and nodes, allowing the transform to maintain continuity of design behavior across distributed execution. By interpreting distributed execution through a single design, the transform collapses incidental complexity and enables deterministic understanding of distributed system behavior.

Lossless Behavioral Representation and Domain-Specific Compression

However, all of this is only possible if the behavior of the design is represented losslessly and faithfully. Any gaps in behavioral representation break the automatic nature of the process, as execution can no longer be understood purely through observation. Instead, missing behavior must be inferred, reintroducing probability and uncertainty into the analysis.


The same requirement applies to testing and learning. If the environments that motivated semantic invariants cannot be unambiguously identified from execution, then testing cannot be automated and constraints cannot be reliably validated. For this reason, it is critical that execution is preserved losslessly.

To make this feasible, the unambiguous nature of the design is leveraged to fully specify the structure of the execution data. Because the design defines exactly what behavior is possible, this structure can be exploited to apply domain-specific compression, significantly reducing the size of the resulting execution trace. Since program execution is highly repetitive and predictable, this compression can be further improved by leveraging recurring structural patterns in the data.

Through this process, the complete execution history and a lossless representation of design behavior can be preserved in a practical and scalable form. To manage these domain-compressed execution traces, an open-source tool named CLP can be used. CLP supports domain-specific compression of unambiguous dynamic traces and enables efficient search directly over compressed data, without requiring decompression. It has been proven at

petabyte scale, making it a practical foundation for enabling automated debugging, testing, and systems management.

The Design Feedback Loop



Through instrumentation, domain-specific compression, and semantic transformation, the entire process is encapsulated in the feedback loop illustrated above. The design serves as the highest authority over the software system, ensuring that execution is always understood with semantic clarity rather. When execution enters a semantically invalid state, the behavior of the design is examined directly to determine whether the cause can be explained by an existing invariant violation. If no such explanation exists, the process enables the design to unambiguously learn from its environment by incorporating new invariants or behavioral refinements.

Through this loop, both the evolution of the design and the realities that shaped it are preserved losslessly. The design does not merely guide implementation; it continuously audits execution, validates its own assumptions, and evolves in response to observed reality, without ever dissolving into ambiguity or probabilistic reasoning.

The Design Learning Platform

Missing (so far):

- In the design section, be more explicit about how failures are predicted and resolved, you just talked about it indirectly.
- In a more abstract way, talk about how the mechanism through which it solves the problem by explaining how it instruments the solution directly.
- Talk about how it transforms systems management by establishing a design learning platform that has clarity from the start while also fully optimizing its implementation
- Conclude by explaining how to establish a fully automated diagnostic tool that automates debugging, testing, systems management and explain how it transforms traditional observability by eliminating the problem it is trying to solve.